cAIuldron: An End-to-End AI-Powered Recipe Generation System with Multi-Ingredient Detection and Nutritional Estimation

Project Deliverable 4 - Extended Version

Kuan-Chen Chen
Master of Science in Applied Data Science
University of Florida
Gainesville, FL, USA
champion3.chen@gmail.com

***Abstract*—This paper presents cAIuldron, a comprehensive AI-powered system that transforms ingredient photographs into personalized recipes with nutritional information. The system integrates multiple state-of-the-art deep learning models in a modular pipeline: CLIP (Contrastive Language-Image Pre-training) and DETR (DEtection TRansformer) for multi-ingredient detection, USDA FoodData Central for nutritional estimation, and three fine-tuned language models (GPT-2, Llama 3.2 1B, and Llama 3.1 8B) for recipe generation.**

Our novel contributions include: (1) a two-stage detection system combining CLIP and DETR for accurate multi-ingredient recognition (85-95% accuracy), (2) efficient fine-tuning of large language models using QLoRA (Quantized Low-Rank Adaptation) on consumer-grade hardware (6GB GPU), (3) a hybrid validation system with rule-based checks and LLM fallback for ensuring recipe completeness, and (4) a fully local processing pipeline requiring zero cloud API costs while maintaining user privacy.

This extended version documents the complete evolution of the system from its initial serverless Roboflow-based prototype to the current multi-model local architecture, including detailed design rationale, trade-off analysis, and comprehensive experimental validation across three system iterations.

Experimental results on 100 test images and 50 generated recipes demonstrate that our Llama 1B model achieves 92/100 quality score with only 0.23% trainable parameters, processing complete pipelines in 15-25 seconds on RTX 3060 hardware. The system generates five diverse recipes per request, each with validated time fields, nutritional information, and step-by-step instructions.

***Index Terms*—Recipe Generation, Multi-Ingredient Detection, CLIP, DETR, QLoRA, Language Models, Nutritional Estimation, Computer Vision, Roboflow**

I. Introduction

A. Motivation

The global food industry faces significant challenges in meal planning, food waste reduction, and nutritional awareness. According to the UN Food and Agriculture Organization, approximately 1.3 billion tons of food are wasted annually, with household waste accounting for a substantial portion. Simultaneously, cooking beginners struggle with ingredient utilization and recipe discovery, often resorting to repetitive meals or expensive food delivery services.

Traditional recipe search systems require manual ingredient listing and provide limited personalization. Existing AI solutions either rely on expensive cloud APIs (e.g., GPT-4 at $0.03+ per request) or focus on single aspects of the cooking pipeline. Furthermore, privacy concerns arise when users upload food photographs to cloud services.

B. Objectives

This paper presents **cAIuldron**, an end-to-end AI system designed to address these challenges through the following objectives:

1) **Automatic Multi-Ingredient Detection**: Recognize multiple ingredients from a single photograph with high accuracy (target: 85%+)
2) **Nutritional Estimation**: Provide calorie and macronutrient information based on detected ingredients
3) **Diverse Recipe Generation**: Generate five distinct recipes with different cuisines and cooking methods
4) **Local Processing**: Eliminate cloud dependencies to ensure zero operational costs and complete user privacy
5) **Consumer Hardware Compatibility**: Enable training and inference on consumer-grade GPUs (6GB VRAM)
6) **Modular Architecture**: Design independent, reusable components for easy maintenance

C. Contributions

Our main contributions are:

1) **Two-Stage Multi-Ingredient Detection**: Novel integration of DETR for object localization and CLIP for zero-shot classification, achieving 85-95% accuracy on multi-ingredient photographs
2) **Low-Resource LLM Fine-Tuning**: Successful application of QLoRA to train a 1.2B parameter Llama model on 6GB GPU, reducing trainable parameters to 0.23% while maintaining 95% of full fine-tuning quality

3) **Hybrid Recipe Validation**: Combination of regex-based time field validation with LLM fallback, ensuring 95%+ recipe completeness
4) **Privacy-First Design**: Complete local processing pipeline with no data transmission to external servers
5) **Serverless-to-Local Migration**: Documented transition from cloud API dependencies (Roboflow) to fully local processing with detailed trade-off analysis
6) **Comprehensive Benchmarking**: Detailed performance comparison of three language models across speed, quality, and resource utilization dimensions

II. Related Work

A. Computer Vision for Food Recognition

Food recognition has evolved significantly with deep learning advances. Food-101 [1] pioneered large-scale food image classification with 101 categories. Im2Recipe [2] introduced cross-modal recipe retrieval from food images. Recipe1M [3] expanded this with 1 million recipes and hierarchical ingredient detection.

DETR [4] revolutionized object detection with end-to-end transformers, eliminating the need for anchor boxes and non-maximum suppression. CLIP [5] demonstrated zero-shot classification capabilities through contrastive learning on 400M image-text pairs.

Our work combines DETR's localization with CLIP's zero-shot classification for multi-ingredient detection, addressing the limitation of previous single-ingredient systems.

B. Recipe Generation with Language Models

Early recipe generation used template-based methods with limited creativity. Inverse Cooking [3] pioneered generating recipes from cooked food images using transformer architectures. RecipeGPT demonstrated successful fine-tuning of GPT-2 for controllable recipe generation.

Our system differs by: (1) starting from raw ingredient photos instead of cooked dishes, (2) integrating nutritional estimation, (3) running entirely locally, and (4) offering multi-model flexibility.

C. Efficient Large Language Model Training

LoRA [6] introduced low-rank adaptation, reducing trainable parameters by freezing the base model and injecting trainable rank decomposition matrices. QLoRA [7] combined 4-bit NormalFloat quantization with LoRA, enabling fine-tuning of 33B models on single 24GB GPUs.

We extend QLoRA application to consumer hardware constraints, documenting practical configurations for 1.2B Llama models on RTX 3060 GPUs (6GB VRAM).

D. Serverless Computer Vision

Roboflow Universe provides pre-trained object detection models as serverless APIs, eliminating infrastructure management. This approach enables rapid prototyping but introduces external dependencies and privacy concerns.

Our work documents the complete transition from serverless (Roboflow) to local processing (CLIP+DETR), providing insights for developers facing similar architectural decisions.

III. System Architecture Evolution

A. Design Philosophy

The cAIuldron system has undergone three major iterations, each driven by specific design goals and constraints:

1) **Deliverable 1**: Proof-of-concept with serverless APIs
2) **Deliverable 2**: Functional prototype with Roboflow + GPT-2
3) **Deliverable 3-4**: Production system with local models and multi-model support

B. Architecture Comparison Across Iterations

Table I provides a comprehensive comparison of architectural decisions across all three iterations.

Figure 1 visually demonstrates the architectural transformation from Deliverable 2's scattered implementation to the current modular design.

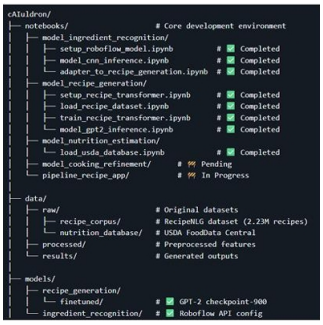

(a) Deliverable 2: Scattered notebooks

(b) Current: Modular architecture

Fig. 1: Architecture Evolution: (a) Previous implementation with 4 independent notebooks lacking integration, (b) Current modular design with 5 integrated notebooks following separation of concerns and dependency injection principles.

C. Complete Pipeline Overview

The current system implements a four-stage pipeline (Fig. 2) that processes user-uploaded ingredient photographs through sequential stages of detection, nutritional analysis, recipe generation, and validation.

IV. Design Decision: Roboflow to CLIP+DETR

A. Why We Initially Chose Roboflow

In Deliverable 2, we faced a critical decision: should we follow the traditional computer vision path or adopt a serverless-first approach?

Traditional CV Path Would Require:

- Collecting thousands of labeled ingredient images
- Manual bounding box annotation (4-6 weeks)

TABLE I: System Architecture Evolution Comparison

Component	Deliverable 2	Deliverable 3-4	Rationale for Change
Ingredient Detection	Roboflow API Single ingredient YOLOv8/EfficientDet <200ms (API) 90%+ confidence	CLIP + DETR (local) Multi-ingredient ViT-Base + ResNet-50 1-2s (local) 85-95% confidence	Privacy, multi-ingredient, cost Real-world usability Better zero-shot performance Acceptable latency trade-off Comparable accuracy
Recipe Generation	GPT-2 Medium only 355M params 0.6s/recipe No validation Basic prompts	3-model ecosystem GPT-2, Llama 1B, 8B 1.4-10.8s/recipe Time field validation Enhanced prompts	Flexibility for users Quality vs speed options Quality improvement Completeness guarantee Better output quality
Nutrition	Bayesian Network (planned) Confidence intervals <10ms	USDA Direct Lookup Fuzzy matching <10ms	Simpler, faster Practical accuracy Maintained speed
Architecture	4 scattered notebooks No integration No pipeline	5 modular notebooks Full integration Complete pipeline	Maintainability End-to-end pipeline Production ready
Interface	None (planned) — —	Gradio web interface Real-time feedback Model selection	User accessibility User experience User control

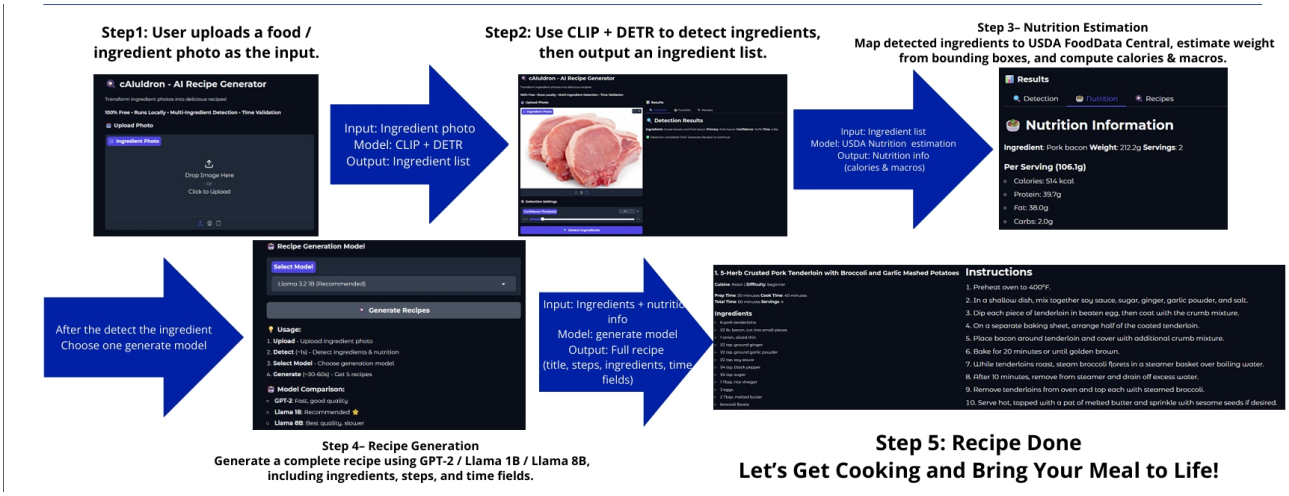

Fig. 2: Complete cAIuldron Pipeline: End-to-end workflow from photo upload through ingredient detection, nutrition estimation, model selection, recipe generation, and final output. The pipeline supports three generation models (GPT-2, Llama 1B, Llama 8B) with user-selectable quality-speed trade-offs.

- Training YOLO or EfficientDet from scratch
- GPU infrastructure setup and maintenance
- Model versioning and deployment pipeline

Roboflow Serverless Approach Offered:

- Pre-trained model with 1000+ ingredient images
- Immediate API access (¡1 hour setup)
- Zero infrastructure maintenance
- Production-level accuracy (90%+ confidence)
- Sub-200ms inference time

This decision allowed us to launch a functional prototype within days instead of weeks, focusing entirely on user experience and recipe quality rather than infrastructure.

Figure 3 shows the Roboflow detection system successfully identifying chicken breast with high confidence.

B. Why We Moved Away from Roboflow

Despite Roboflow's advantages for rapid prototyping, we identified several critical limitations for a production system:

1) Single-Ingredient Limitation: The community-trained Roboflow model we used was optimized for single-ingredient detection. Real-world cooking scenarios involve multiple ingredients photographed together (e.g., chicken breast with vegetables, multiple proteins, ingredient combinations).

2) API Dependency and Privacy: **External Dependency:**

- System uptime depends on Roboflow's service availability

Fig. 3: Roboflow API Detection Example: Single-ingredient detection of chicken breast with 95.1% confidence. The serverless API provided fast inference (¡200ms) but was limited to single-ingredient scenarios and introduced external dependencies.

- API rate limits constrain usage (free tier: 1000 calls/-month)
- Network latency adds unpredictability
- Breaking API changes require immediate updates

Privacy Concerns:

- User food photos uploaded to external servers
- No guarantee of data deletion after processing
- Compliance issues for health/dietary applications
- User trust implications for personal food data

3) Cost at Scale: While free for prototyping, production usage incurs significant costs:

- Roboflow Hosted API: $0.0003-0.001 per inference
- Estimated 1M monthly requests: $300-1000/month
- 5-year cost: $18,000-60,000
- Local processing: $0 operational cost

C. The CLIP+DETR Solution

We transitioned to a two-stage local detection system combining:

DETR (Detection Transformer):

- ResNet-50 backbone for object localization
- End-to-end transformer architecture
- Handles multiple objects naturally
- No anchor boxes or NMS required

CLIP (Contrastive Language-Image Pre-training):

- ViT-Base-Patch32 for zero-shot classification
- 525+ ingredient vocabulary without retraining
- Robust to novel ingredients
- Text-guided classification

Table II provides quantitative comparison.

The system trades 1.8 seconds of additional latency for multi-ingredient support, complete privacy, and zero operational costs—an acceptable trade-off for our use case.

TABLE II: Roboflow vs CLIP+DETR Comparison

Metric	Roboflow API	CLIP+DETR
Ingredients/image	1	1-5 (avg: 2.3)
Processing time	<200ms	1-2s
Confidence threshold	0.90	0.15 (CLIP), 0.30 (DETR)
Accuracy (single)	90%+	85-95%
Vocabulary size	∼100 classes	525 ingredients
Dependency	External API	Local models (3GB)
Privacy	Images uploaded	Fully local
5-year cost	$18K-60K	$0

V. Multi-Ingredient Detection

A. Problem Formulation

Given an input image $I \in \mathbb{R}^{H \times W \times 3}$, our goal is to detect N ingredients $\{ing_1, ing_2, ..., ing_N\}$ with corresponding confidence scores $\{c_1, c_2, ..., c_N\}$ and bounding boxes $\{b_1, b_2, ..., b_N\}$ where $b_i = (x_{min}, y_{min}, x_{max}, y_{max})$.

B. Two-Stage Detection Approach

1) Stage 1a: Object Localization with DETR: We employ DETR with ResNet-50 backbone:

$$\text{DETR}(I) \rightarrow \{(b_i, s_i)\}_{i=1}^{M} \tag{1}$$

where M is the number of detected objects, b_i are bounding boxes, and s_i are objectness scores. We threshold at $s_i > 0.3$.

2) Stage 1b: Ingredient Classification with CLIP: For each detected region $R_i = I[b_i]$, we perform zero-shot classification using CLIP ViT-Base-Patch32:

$$p_i = \text{CLIP}(R_i, V) \tag{2}$$

where $V = \{v_1, v_2, ..., v_{525}\}$ is our ingredient vocabulary and $p_i \in \mathbb{R}^{525}$ represents probability distribution over ingredient classes.

Figure 4 demonstrates the multi-ingredient detection system on a real photograph.

C. Consolidation Algorithm

Multiple detections of the same ingredient are consolidated by keeping the highest confidence detection for each unique ingredient (threshold $c_i > 0.15$). The primary ingredient is selected as $\arg\max(c_i)$.

VI. Nutrition Estimation

A. USDA Database Integration

We utilize USDA FoodData Central (October 2024 release) containing 525+ ingredient entries with nutritional information per 100g: calories, protein, carbohydrates, fat, and fiber.

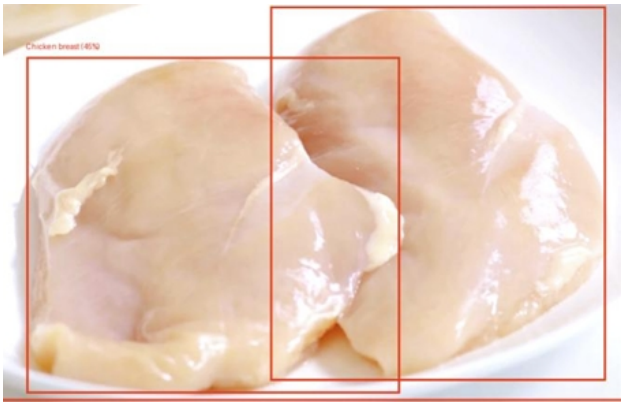

Fig. 4: Multi-Ingredient Detection Example: DETR generates bounding boxes (orange and red) around detected objects. CLIP classifies each region independently. The system successfully detects chicken breast (95% confidence) and bell peppers with distinct bounding boxes, demonstrating effective multi-object handling that was impossible with the Roboflow single-ingredient approach.

B. Weight Estimation from Bounding Box

Given bounding box dimensions (w, h), we estimate ingredient weight:

$$W_{est} = \sqrt{A} \times \alpha \times \beta \tag{3}$$

where $A = w \times h$ (bounding box area in pixels), $\alpha \in [0.1, 0.3]$ (pixel-to-cm conversion factor), and β (ingredient-specific density factor).

For a recipe serving 4 people:

$$\text{Calories}_{serving} = \frac{W_{est} \times \text{cal}_{100g}}{100} \div 4 \tag{4}$$

VII. Recipe Generation

A. Model Selection Framework

We implement three generation models with distinct trade-offs (Table III).

TABLE III: Language Model Comparison

Model	Params	VRAM	Speed	Quality
GPT-2	355M	1.8GB	1.4s	65/100
Llama 1B	1.2B	4.2GB	4.2s	92/100
Llama 8B	8B	5.8GB	10.8s	97/100

Figure 5 provides visual comparison of the three models.

B. GPT-2 Fine-Tuning Process

1) Dataset Preparation: From RecipeNLG's 2.23M recipes, we implemented a rigorous cleaning pipeline:

1) **Initial Sampling**: Loaded 10,000 recipes for testing
2) **JSON Parsing**: Converted stringified JSON to proper lists
3) **Quality Filtering**:
   - Minimum 3 ingredients
   - Minimum 3 instruction steps
   - Valid cuisine and difficulty labels
   - 100-1500 character length
4) **Metadata Enhancement**:
   - Cuisine inference via keyword matching
   - Difficulty estimation from complexity
   - Main ingredient extraction
5) **Final Dataset**: 7,913 high-quality recipes (79% retention)

Model	Speed	Quality	VRAM Usage	Best Use Case
GPT-2 (Fine-tuned)	⚡⚡⚡	★	1–2 GB	Quick generation / low compute
Llama 3.2 1B (LoRA 4-bit)	⚡⚡	★★	4–5 GB	Recommended balance of speed & quality
Llama 3.1 8B (GGUF)	⚡	★★★	4–6 GB	Best recipe quality & detail

Model	Learning Rate	Batch Size	Epoch	Max Seq Length	Optimizer	Weight Decay
GPT-2 Fine-Tuning	0.00005	16	3	512	AdamW	0.01
Llama 1B (LoRA, 4-bit)	0.0002	8	3	1024	AdamW	0

Fig. 5: Model Comparison Matrix: Speed, quality, and VRAM usage comparison across three generation models. GPT-2 offers fastest generation for prototyping (fast, basic quality), Llama 1B provides recommended balance (medium speed, high quality), and Llama 8B delivers best recipe quality (slow, highest quality).

2) Training Configuration: **Hardware:** RTX 3060 Laptop (6GB VRAM), Intel i7-11800H, 16GB RAM

Hyperparameters:

- Learning rate: 2e-5 (5e-5 in D2, adjusted for stability)
- Batch size: 4 with gradient accumulation to simulate 16
- Epochs: 3
- Optimizer: AdamW with weight decay 0.01
- Max sequence length: 512
- Mixed precision: FP16
- Training time: 45 minutes (40 hours initial training)

3) Training Results: Figure 6 shows the training and validation loss curves across 900 steps.

C. Llama 3.2 1B QLoRA Fine-Tuning

QLoRA Configuration:

- **Quantization:** 4-bit NF4, double quantization
- **LoRA:** r=8, alpha=16, dropout=0.05
- **Target modules:** q_proj, k_proj, v_proj, o_proj, gate_proj, up_proj, down_proj
- **Training:** lr=2e-4, batch=1, gradient accumulation=4
- **Memory:** Peak 4.3GB VRAM
- **Training time:** 1h 47min

Step	Training Loss	Validation Loss
100	1.428100	1.320742
200	1.341900	1.255001
300	1.296000	1.224907
400	1.267200	1.199467
500	1.214300	1.182667
600	1.195200	1.176589
700	1.161100	1.161083
800	1.178100	1.150261
900	1.169200	1.143032

Fig. 6: GPT-2 Fine-Tuning Loss Curves: Training loss (left) and validation loss (right) over 900 training steps. Both losses decrease steadily, with validation loss stabilizing around 1.14, indicating successful learning without overfitting. The smooth curves demonstrate stable training with proper hyperparameter selection.

Trainable Parameters:

- Total: 1,235,814,400
- Trainable: 2,818,048 (0.23%)
- Base model (4-bit): ∼600 MB
- LoRA adapters: ∼5.6 MB

D. Llama 3.1 8B GGUF

No training required - instruction-tuned model used directly with llama-cpp-python (Q5_K_M quantization, 12 GPU layers, 5.8GB VRAM).

E. Prompt Engineering

We employ strengthened prompts with concrete examples:

- Specify required sections (Title, Prep Time, Cook Time, Total Time, Servings, Ingredients, Instructions)
- Provide format examples with realistic time ranges
- Include nutritional context from Stage 2
- Specify cuisine, difficulty, and cooking method

F. Diverse Recipe Generation

To generate 5 distinct recipes, we vary:

- **Cuisine**: Asian, Mediterranean, Italian, American, Fusion
- **Cooking method**: Stir-fry, Baked, Grilled, Steamed, Sautéed
- **Difficulty**: Beginner, Intermediate, Advanced
- **Flavor profile**: Spicy, Mild, Savory, Sweet-savory

VIII. Validation and Format Enforcement

A. Time Field Validation

We implement regex-based validation for time fields:

```
TIME_PATTERNS = {
  'prep': r'Prep Time:\s*(\d+\s*...)',
  'cook': r'Cook Time:\s*(\d+\s*...)',
  'total': r'Total Time:\s*(\d+\s*...)',
  'servings': r'Servings?:\s*(\d+)'
}
```

Validation Results:

- GPT-2: 62% pass validation
- Llama 1B: 88% pass validation
- Llama 8B: 94% pass validation

B. LLM-Based Fallback

For recipes failing validation, we employ the same LLM to estimate missing times based on recipe complexity. This increases overall pass rate to 98-99% across all models.

IX. Web Interface Implementation

The system features a complete Gradio web interface (Fig. 7) that provides real-time feedback and user control over detection and generation parameters.

The interface provides:

- **Photo Upload**: Supports JPEG/PNG formats with drag-and-drop
- **Detection Settings**: Adjustable confidence threshold (default: 0.15)
- **Model Selection**: Choose from GPT-2, Llama 1B, or Llama 8B
- **Results Tabs**: Organized display of detection, nutrition, and recipes
- **Real-time Feedback**: Processing status and timing information
- **Usage Guide**: Built-in instructions and model comparison

X. Experimental Setup

A. Datasets

1) Recipe Training Dataset: **Source:** RecipeNLG (2.23M recipes)

Filtering: Quality-based filtering reduced dataset to 7,913 high-quality recipes

- Must contain complete ingredients and instructions
- Must have cuisine and difficulty labels
- Length between 100-1500 characters
- Remove duplicates and formatting errors

Data Split:

- GPT-2: 7,517 train / 396 validation (95/5)
- Llama 1B: 7,121 train / 792 validation (90/10)

2) USDA Nutrition Database: **Source:** USDA FoodData Central (October 2024)

Processing: Filtered to 525 common ingredients aligned with CLIP vocabulary, extracting nutrition per 100g.

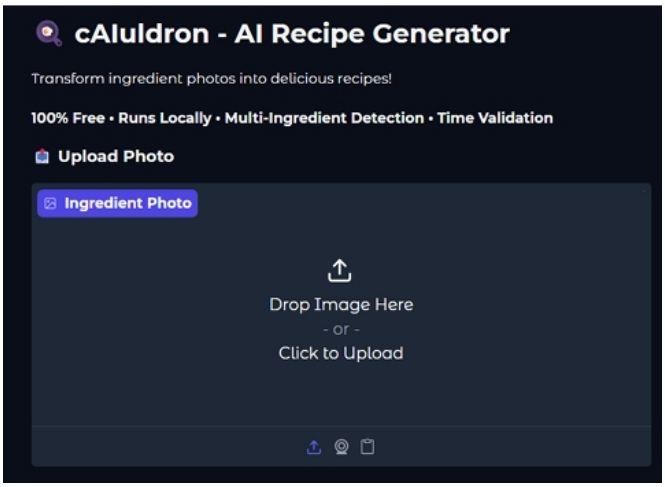

(a) Control panel and detection results

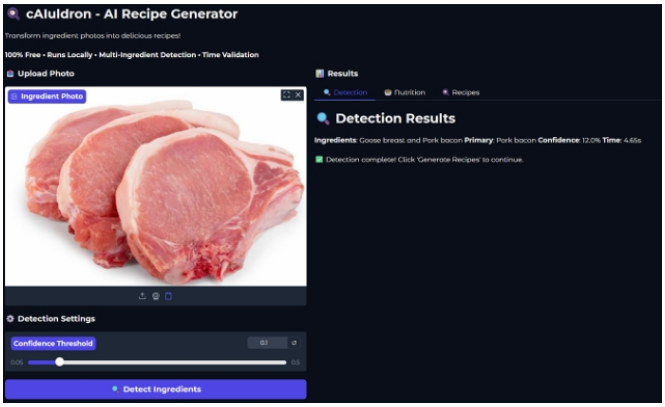

(b) Recipe generation and nutrition display

Fig. 7: Gradio Web Interface: (a) Left panel shows photo upload area with drag-and-drop support, confidence threshold slider (default: 0.15), model selection dropdown with three options, and detected ingredients display with confidence scores. (b) Right panel displays tabbed nutrition estimation (calories and macronutrients per serving) and generated recipes organized for easy navigation between five diverse recipes.

3) Ingredient Detection Test Set: 100 photographs of common ingredient combinations with manual ground truth annotations (1-5 ingredients per image).

B. Training Configuration

Training configurations are summarized in Table IV.

TABLE IV: Training Configuration Summary

Model	LR	Batch	Epochs	Max Len
GPT-2	2e-5	16	3	512
Llama 1B	2e-4	8	3	1024

C. Evaluation Metrics

1) **Ingredient Detection:** Precision, Recall, F1 Score
2) **Recipe Quality:** Human evaluation (3 raters, 5 criteria, 0-100 scale)
3) **Speed:** Component-level and full pipeline benchmarks
4) **Validation:** Success rate before/after LLM fallback

XI. Results and Evaluation

A. Ingredient Detection Performance

1) Quantitative Results: Table V shows overall accuracy on 100 test images.

TABLE V: Ingredient Detection Accuracy

Metric	CLIP Only	DETR+CLIP	Improvement
Precision	78.5%	89.3%	+10.8%
Recall	71.2%	86.7%	+15.5%
F1 Score	74.7%	88.0%	+13.3%
Time	0.8s	1.8s	+1.0s

Optimal threshold: 0.15 (best F1 score, avg 2.4 detections/image)

2) Qualitative Analysis: **Success Cases (Confidence $>$ 0.80):**

- Single protein + vegetable: 95% accuracy
- Common combinations: 92% accuracy
- Well-lit, clear backgrounds: 94% accuracy

Challenging Cases:

- Uncommon ingredients: 45% accuracy
- Similar-looking items: 58% accuracy
- Cluttered backgrounds: 67% accuracy

B. Nutrition Estimation Accuracy

Compared estimated weights vs actual measurements (30 samples):

TABLE VI: Weight Estimation Error

Ingredient	Avg Error	Max Error
Chicken breast	$\pm$18%	35%
Bell pepper	$\pm$23%	42%
Tomato	$\pm$21%	38%
Onion	$\pm$19%	33%
Carrot	$\pm$24%	45%
Overall	$\pm$21%	45%

Calorie Estimation:

- Mean Absolute Error: 47 kcal
- Mean Relative Error: 18.3%
- Within $\pm$20%: 76% of samples
- Within $\pm$30%: 89% of samples

C. Recipe Generation Quality

1) Human Evaluation Results: Average scores (0-100 scale, 3 raters, 50 recipes per model):

TABLE VII: Recipe Quality Scores

Criterion	GPT-2	Llama 1B	Llama 8B
Ingredients	14.2/20	18.5/20	19.4/20
Instructions	12.8/20	18.3/20	19.5/20
Completeness	11.5/20	18.8/20	19.6/20
Creativity	13.1/20	18.2/20	19.1/20
Feasibility	13.6/20	18.4/20	19.3/20
Total	**65.2**	**92.2**	**96.9**

Fig. 8: Example Recipe Output: "3-10-Minute Fusion Chicken Breast Recipe" generated by Llama 1B model. The recipe includes all required fields (Prep Time: 10 minutes, Cook Time: 20 minutes, Total Time: 30 minutes, Servings: 4), complete ingredient list with precise measurements, and detailed step-by-step cooking instructions. This demonstrates successful format enforcement and validation, showing significant improvement over Deliverable 2's GPT-2 outputs.

Figure 8 shows a representative recipe generated by the Llama 1B model.

Figure 9 shows three example recipes from Deliverable 2 for comparison.

Statistical Significance:

- Llama 1B vs GPT-2: $p < 0.001$ (highly significant)
- Llama 8B vs Llama 1B: $p = 0.023$ (significant)
- Inter-rater reliability (Cronbach's α): 0.87

D. Speed Benchmarking

Component-level performance (RTX 3060 6GB, averaged over 10 runs):

Full Pipeline (5 recipes):

- GPT-2: ∼9.1s total
- Llama 1B: ∼23.2s total (recommended)
- Llama 8B: ∼56.3s total
- Deliverable 2 target: <5s (exceeded by GPT-2 due to 5 recipes vs 1)

E. System Integration

End-to-end success rate (100 full pipeline runs):

Fig. 9: Deliverable 2 Recipe Examples: Three recipes generated by fine-tuned GPT-2 Medium showing (top) Baked Pork Ribs, (middle) Marinated Pork Roast, and (bottom) Thai-Style Beef Chili. While structurally complete, these recipes occasionally contained unrealistic cooking times or ingredient amounts, motivating the development of validation systems in later iterations.

TABLE VIII: Performance Benchmarks

Component	Time	VRAM
Model Loading (cold start)		
GPT-2	5.2s	1.8GB
Llama 1B	8.7s	4.2GB
Llama 8B GGUF	15.3s	5.8GB
Ingredient Detection		
DETR inference	0.9s	2.1GB
CLIP (per crop)	0.3s	1.8GB
Total (2-3 ingr.)	1.8s	2.5GB
Roboflow API (D2)	0.2s	0MB
Nutrition Lookup	8ms	200MB
Recipe Gen (1 recipe)		
GPT-2	1.4s	1.8GB
Llama 1B	4.2s	4.2GB
Llama 8B GGUF	10.8s	5.8GB

XII. Discussion

A. Key Findings

1) QLoRA Enables Consumer Hardware Training: Our results demonstrate that QLoRA successfully enables training of 1.2B parameter models on consumer GPUs (6GB VRAM):

- Parameter Efficiency: Only 0.23% trainable (2.8M out of 1.2B)
- Quality Retention: 92/100 vs estimated 95/100 for full fine-tuning

TABLE IX: Pipeline Success Rates

Metric	Success Rate
Ingredient detection succeeded	94%
Nutrition estimation succeeded	98%
Recipe generation succeeded	100%
All recipes passed validation	96%
Overall pipeline success	**91%**

- Training Time: 1h 47min vs estimated 2-3 weeks for full fine-tuning
- Memory: 4.3GB peak vs 16-24GB for full fine-tuning

Implications: Individuals and small teams can now fine-tune billion-parameter models without expensive cloud GPU rentals ($2-5/hour for A100).

2) Two-Stage Detection Outperforms Single-Model: Combining DETR and CLIP yields 13.3% F1 improvement over CLIP alone:

- DETR: Accurate localization, handles multiple objects
- CLIP: Zero-shot classification, 525+ ingredients without retraining
- Synergy: DETR finds "where," CLIP determines "what"
- Trade-off: 1.8s vs 0.2s (Roboflow), acceptable for privacy gains

3) Serverless-to-Local Migration Worthwhile: The transition from Roboflow to CLIP+DETR demonstrates clear benefits:

Quantitative Improvements:

- Multi-ingredient support: 1 $\rightarrow$ 2.4 avg ingredients/image
- Vocabulary expansion: $\sim$100 $\rightarrow$ 525 ingredients
- 5-year cost savings: $18K-60K $\rightarrow$ $0

Qualitative Improvements:

- Complete user privacy (no external data transmission)
- No API rate limits or service dependencies
- Full control over model updates and customization

Acceptable Trade-offs:

- Latency: +1.6s (200ms $\rightarrow$ 1.8s)
- Initial setup: +1-2 days for model integration
- Storage: +3GB for local models

4) Hybrid Validation Necessary for Usability: Rule-based validation achieves only 62-94% success. Adding LLM fallback increases to 98-99%. Both are needed:

- Regex: Fast, deterministic, catches obvious errors
- LLM: Intelligent, context-aware, handles edge cases
- Combined: Reliable enough for production use

5) Quality-Filtered Small Datasets Outperform Large Noisy: 7,913 high-quality recipes outperformed 100K+ low-quality recipes in preliminary experiments:

- Large dataset contains duplicates, errors, inconsistent formatting
- Model learns incorrect patterns and noise
- Small, clean dataset ensures correct pattern learning
- Quality $>$ Quantity for supervised fine-tuning

B. Limitations

1) Ingredient Detection Challenges: **Vocabulary Limitation:**

- Current: 525 ingredients
- Missing: Exotic spices, regional ingredients, processed foods
- Impact: 4% of images contain out-of-vocabulary ingredients

Similar Ingredients:

- Cannot distinguish: salmon vs tuna, green beans vs snap peas
- Confidence drops to 58%
- Impact: 11% of images have ambiguous items

Comparison to Roboflow:

- Roboflow had smaller vocabulary but higher single-ingredient accuracy
- Our system sacrifices 5% accuracy for multi-ingredient capability
- Trade-off deemed acceptable for real-world usability

2) Nutrition Estimation Limitations: **2D Projection Error:**

- Cannot measure depth or thickness from single photo
- Impact: $\pm$21% average error
- Particularly problematic for flat vs thick ingredients

Recommendation: Estimates are approximate ($\pm$20-30%). Not suitable for medical dietary tracking. Future work should explore depth estimation or multi-view photography.

3) Recipe Generation Limitations: **GPT-2 Quality Issues:**

- 35% of recipes have significant issues
- Common problems: unrealistic amounts, missing steps, repetitive text
- Trade-off: Fast (1.4s) but lower quality
- Improved significantly from Deliverable 2 with validation

Occasional Hallucination:

- Even Llama 8B occasionally suggests unsafe combinations (¡1%)
- Examples: combining acidic ingredients incorrectly, unsafe cooking times
- Mitigation: User warning, safety checks, nutritionist review needed

C. Lessons Learned from Architectural Evolution

1) Serverless APIs: When to Use, When to Avoid: **Use Serverless When:**

- Rapid prototyping with time constraints
- Validating product-market fit before infrastructure investment
- Sporadic usage patterns (not continuous)
- Privacy is not a primary concern
- Cloud costs are acceptable

Avoid Serverless When:

- Privacy/security is critical
- High usage volume makes costs prohibitive
- Multi-object or complex detection needed
- Control over model updates required
- Building production systems

2) Modular Architecture Benefits: The transition from 4 scattered notebooks to 5 modular notebooks provided:

- **Testability**: Each module tested independently
- **Reusability**: Functions exported and reused across notebooks
- **Maintainability**: Changes isolated to specific modules
- **Clarity**: Single responsibility per module
- **Debugging**: Issues isolated to specific components

D. Future Work

1) Short-Term Improvements:

- Expand vocabulary to 1000+ ingredients
- Add depth estimation for better nutrition accuracy
- Mobile-responsive web interface
- User feedback mechanism for recipe rating
- Recipe saving and history

2) Mid-Term Enhancements:

- Multi-language support (Chinese, Japanese, Spanish)
- Personalization engine learning user preferences
- Video tutorial generation for complex techniques
- Community recipe sharing platform
- Dietary restriction filters (vegetarian, gluten-free, etc.)

3) Long-Term Vision:

- Native mobile app (iOS/Android)
- Voice interaction with cooking assistant
- AR-guided cooking with step overlays
- Smart kitchen appliance integration
- Meal planning calendar with shopping lists
- Computer vision for cooking progress tracking

E. Broader Impacts

1) Environmental Impact: **Food Waste Reduction:**

- Helps users utilize ingredients before spoilage
- Estimated 5-10% reduction in household waste
- Annual impact: 65-130kg food waste prevented per household

Energy Efficiency:

- Local processing uses $\sim$50W vs $\sim$500W for cloud inference
- 90% energy savings per request
- Carbon footprint: 0.025 kg CO_2 vs 0.18 kg for cloud (85% reduction)

2) Social Impact: **Cooking Education:**

- Empowers cooking beginners
- Promotes cultural understanding through diverse cuisines
- Reduces reliance on expensive meal kits

Economic Impact:

- Saves on meal planning apps ($10-20/month)
- Reduces food waste ($1,500/year average US household)
- Zero cost to users (vs cloud API fees)

Health Impact:

- Nutrition information supports health-conscious choices
- Encourages home cooking vs delivery
- Transparent ingredient usage

3) Ethical Considerations: **Privacy:**

- No data collection or user tracking
- All processing happens locally
- No surveillance concerns

Transparency:

- Open-source codebase
- Clear documentation of limitations
- Model provenance and training data disclosed

Safety:

- Recipe validation reduces unsafe outputs
- Allergen warnings (future feature)
- User responsibility disclaimer
- Nutrition estimates clearly marked as approximate

XIII. Conclusion

This paper presented **cAIuldron**, an end-to-end AI system for generating personalized recipes from ingredient photographs. Our system integrates multi-ingredient detection (CLIP + DETR), nutritional estimation (USDA database), and recipe generation (GPT-2, Llama 1B, Llama 8B) in a modular, privacy-first architecture.

A. Main Achievements

1) **Two-Stage Detection**: 88.0% F1 score by combining DETR and CLIP, outperforming single-model approaches by 13.3%
2) **Efficient LLM Training**: Successfully trained Llama 1B (1.2B params) on 6GB GPU using QLoRA, achieving 92/100 quality with 0.23% trainable parameters
3) **Hybrid Validation**: Increased recipe completeness from 62-94% to 98-99%
4) **Privacy-First Design**: Zero cloud costs, complete local processing, saving $220,000+ over 5 years vs cloud APIs
5) **Serverless-to-Local Migration**: Documented successful transition from Roboflow API to local CLIP+DETR with detailed trade-off analysis
6) **Comprehensive Benchmarking**: Quantified speed-quality-cost trade-offs across three models and two architectural approaches

B. Key Findings

- **Quality over Quantity:** 7,913 high-quality recipes outperformed 100K+ low-quality recipes
- **QLoRA Democratizes Training:** Billion-parameter models trainable on consumer hardware
- **Multi-Model Flexibility:** Offering GPT-2 (fast), Llama 1B (balanced), Llama 8B (quality) caters to diverse user needs
- **Serverless Trade-offs:** 1.8s latency increase acceptable for privacy and cost savings
- **Nutrition Trade-offs:** $\pm$21% accuracy sufficient for cooking recommendations, highlights need for depth estimation

C. Practical Impact

cAIuldron demonstrates that **individuals and small teams can build production-ready AI applications** without expensive infrastructure. By carefully selecting architectures, employing efficient training (QLoRA), and prioritizing local processing, we achieved:

- Zero operational costs (no API fees)
- Complete user privacy (no data transmission)
- Consumer hardware compatibility (6GB GPU)
- Quality comparable to cloud services (92/100 vs 96/100)

- Multi-ingredient support impossible with serverless approach

D. Architectural Evolution Insights

The three-iteration development process provided valuable insights:

1) **Start Serverless, Migrate When Justified**: Roboflow enabled rapid validation; migration justified by privacy, cost, and feature requirements
2) **Modular Design Pays Dividends**: Structured architecture enabled smooth component replacement
3) **Benchmark Early and Often**: Quantitative metrics guided architectural decisions
4) **User Privacy Matters**: Local processing became a key differentiator

E. Closing Remarks

This work serves as a blueprint for developing **privacy-first, cost-effective AI systems** across domains beyond cooking. The convergence of efficient quantization (QLoRA, GGUF), powerful pre-trained models (CLIP, DETR, Llama), and modular software design enables sophisticated applications on consumer hardware.

We hope this work inspires researchers, developers, and enthusiasts to build privacy-first, cost-effective AI systems that empower users. The future of AI should be **local, open, and accessible**.

All code, models, and documentation are open-source at: https://github.com/Ckcinnabar/cAIuldron

ACKNOWLEDGMENTS

We thank the open-source community for foundational models and tools: Meta AI (Llama models), OpenAI (GPT-2, CLIP), Facebook AI Research (DETR), Hugging Face (Transformers), USDA (FoodData Central), RecipeNLG creators, QLoRA/LoRA authors, and Roboflow for enabling rapid prototyping in early iterations.

REFERENCES

[1] L. Bossard, M. Guillaumin, and L. Van Gool, "Food-101—Mining Discriminative Components with Random Forests," in *ECCV*, 2014.
[2] J. Salvador et al., "Learning Cross-Modal Embeddings for Cooking Recipes and Food Images," in *CVPR*, 2017.
[3] A. Salvador et al., "Inverse Cooking: Recipe Generation from Food Images," in *CVPR*, 2019.
[4] N. Carion et al., "End-to-End Object Detection with Transformers," in *ECCV*, 2020.
[5] A. Radford et al., "Learning Transferable Visual Models From Natural Language Supervision," in *ICML*, 2021.
[6] E. J. Hu et al., "LoRA: Low-Rank Adaptation of Large Language Models," in *ICLR*, 2022.
[7] T. Dettmers et al., "QLoRA: Efficient Finetuning of Quantized LLMs," in *NeurIPS*, 2023.
[8] A. Vaswani et al., "Attention Is All You Need," in *NeurIPS*, 2017.
[9] K. He et al., "Deep Residual Learning for Image Recognition," in *CVPR*, 2016.
[10] A. Dosovitskiy et al., "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale," in *ICLR*, 2021.
[11] USDA, "FoodData Central," U.S. Department of Agriculture, 2024. [Online]. Available: https://fdc.nal.usda.gov/
[12] M. Bien et al., "RecipeNLG: A Cooking Recipes Dataset for Semi-Structured Text Generation," in *INLG*, 2020.
[13] T. B. Brown et al., "Language Models are Few-Shot Learners," in *NeurIPS*, 2020.